

# Stochastic gradient descent

Victor Kitov

[v.v.kitov@yandex.ru](mailto:v.v.kitov@yandex.ru)

# Table of Contents

- 1 Gradient properties reminder
- 2 Gradient descent optimization

# Gradient

- For any function  $f(x)$ , depending from  $x = (x_1, \dots, x_D)^T$  gradient

$$\nabla f(x) := \begin{pmatrix} \frac{\partial f(x)}{\partial x_1} \\ \frac{\partial f(x)}{\partial x_2} \\ \dots \\ \frac{\partial f(x)}{\partial x_D} \end{pmatrix}$$

- If function  $f(x, y)$  depends on other variables  $y$  gradient  $\nabla_x$  considers only derivatives with respect to  $x$ :

$$\nabla_x f(x, y) := \begin{pmatrix} \frac{\partial f(x)}{\partial x_1} \\ \frac{\partial f(x)}{\partial x_2} \\ \dots \\ \frac{\partial f(x)}{\partial x_D} \end{pmatrix}$$

## Directional derivative

### Definition 1

Consider differentiable function  $f : \mathbb{R}^D \rightarrow \mathbb{R}$ . A derivative along direction  $d$ ,  $\|d\| = 1$  is defined as

$$f'(x, d) = \lim_{\lambda \rightarrow 0} \frac{f(x + \lambda d) - f(x)}{\lambda}$$

### Theorem 2

$$f'(x, d) = \nabla f(x)^T d$$

*Proof.* Using 1-st order Taylor expansion we have

$$\begin{aligned} f(x + \lambda d) &= f(x) + \nabla f(x)^T (\lambda d) + o(\lambda) \\ \frac{f(x + \lambda d) - f(x)}{\lambda} &= \nabla f(x)^T d + o(1) \xrightarrow{\lambda \rightarrow 0} \nabla f(x)^T d \end{aligned}$$



## Direction of maximal growth/decrease

### Theorem 3

*For differentiable function  $f(x)$  locally at point  $x$ :*

- $\frac{\nabla f(x)}{\|\nabla f(x)\|}$  *is the direction of maximum growth*
- $-\frac{\nabla f(x)}{\|\nabla f(x)\|}$  *is the direction of maximal decrease.*

*Proof.* 1-st order Taylor expansion

$$f(x + \lambda d) = f(x) + \nabla f(x)^T (\lambda d) + o(\lambda)$$

From Cauchi-Schwartz inequality, taking  $\|d\| = 1$ :

$$\left| \nabla f(x)^T d \right| \leq \|\nabla f(x)\| \|d\| = \|\nabla f(x)\|$$

Equality is achieved when  $d \propto \nabla f(x)$ , i.e.  
 $d = \pm \nabla f(x) / \|\nabla f(x)\|$ .



# Table of Contents

- 1 Gradient properties reminder
- 2 Gradient descent optimization

# Comments

- Empirical risk minimization


$$L(w) = \frac{1}{N} \sum_{n=1}^N \mathcal{L}(x_n, y_n, w) \rightarrow \min_w$$

- Regression:  $\mathcal{L}(f_w(x_n) - y_n)$
- Classification  $\mathcal{L}((g_{y_n, w}(x_n) - g_{-y_n, w}(x_n)) y_n) = \mathcal{L}(g_w(x_n) y_n)$ .
- Problems:
  - for general  $\mathcal{L}, f(\cdot), g(\cdot)$  no analytical solution
  - $\hat{\beta} = (X^T X)^{-1} X^T Y$  - complexity  $O(D^3)$  - high for big  $D$ .

# Gradient descend optimization

- Gradient descend - iterative movement in steepest descent:

$$w := w - \nabla_w L(w)$$



- If  $\mathcal{L}(u)$ -convex  $\Rightarrow L(w)$ -convex  $\Rightarrow$  local optimum is global optimum.



# Gradient descend optimization

**INPUT:**

- \*  $\varepsilon$ : parameter, controlling the speed of convergence
- \* stopping rule

**ALGORITHM:**

initialize  $t = 0$ ,  $w_0$  randomly

**WHILE** stopping rule is not satisfied:

$$w_{t+1} := w_t - \varepsilon \nabla_w L(w_t)$$


$$t := t + 1$$

**RETURN**  $w_n$


Stopping rules:  $|L(w_t) - L(w_{t-1})|$  or  $\|w_t - w_{t-1}\|$  below threshold,  
or fixed #[iterations].

# Learning rate selection<sup>1</sup>

$\varepsilon$  should be selected carefully based on  $L(w_t)$  dynamics.



Large learning rate: Overshooting.




Small learning rate: Many iterations until convergence and trapping in local minima.

---

<sup>1</sup>Picture [source](#).




# Learning rate selection



# Feature normalization

Convergence is faster for normalized features:

- feature normalization solves the problem of «elongated valleys»



## Problem of gradient descend (GD)

**INPUT:**

- \*  $\varepsilon_t$ : controls the speed of convergence
- \* stopping rule

**ALGORITHM:**

initialize  $t = 0$ ,  $w_0$  randomly

**WHILE** stopping rule is not satisfied:

$$w_{t+1} := w_t - \varepsilon_t \frac{1}{N} \sum_{i=1}^N \nabla_w \mathcal{L}(x_i, y_i | w_n)$$

$$t := t + 1$$

**RETURN**  $w_n$

Gradient calculation requires  $O(N)$  operations!

# Stochastic gradient descent (SGD)

**INPUT:**

- \*  $\varepsilon_t$ : controls the speed of convergence
- \* stopping rule

**ALGORITHM:**

initialize  $t = 0$ ,  $w_0$  randomly

**WHILE** stopping rule is not satisfied:

    randomly sample  $I = \{n_1, \dots, n_K\}$  from  $\{1, 2, \dots, N\}$

$w_{t+1} := w_t - \varepsilon_t \frac{1}{K} \sum_{n \in I} \nabla_w \mathcal{L}(x_n, y_n | w_t)$

$t := t + 1$


**RETURN**  $w_t$

Main idea:  $\frac{1}{N} \sum_{n=1}^N \mathcal{L}(x_n, y_n | w) \approx \frac{1}{K} \sum_{n \in I} \mathcal{L}(x_n, y_n | w)$ , one step takes  $O(K)$ ,  $K \ll N$ ,  $K$ -minibatch size.

## SGD comments

- Indices generation: before each pass through the training set, it is randomly shuffled and then passed sequentially.
- Works even for  $K = 1$ .
- $\frac{1}{K} \sum_{i \in I} \nabla_w \mathcal{L}(x_i, y_i | w_n)$  can be computed in  $O(1)$  for small  $K$  because processors internally perform vector arithmetics.

# Learning rate selection



- Convergence requirements:

$$\sum_t \varepsilon_t = +\infty \quad \text{SGD should reach any point}$$

$$\sum_t \varepsilon_t^2 < +\infty \quad \varepsilon_t \text{ should converge to 0 fast}$$

In practice  $\varepsilon_t = \frac{\alpha}{t+\beta}$  or constant which is reduced when criterion



# SGD reformulated

**INPUT:**

- \*  $\varepsilon_t$ : controls the speed of convergence
- \* stopping rule

**ALGORITHM:**

initialize  $t = 0$ ,  $w_0$  randomly,  $\Delta w_0 = 0$

**WHILE** stopping rule is not satisfied:

    randomly sample  $I = \{n_1, \dots, n_K\}$  from  $\{1, 2, \dots, N\}$

$$\Delta w_{t+1} = -\frac{1}{K} \sum_{n \in I} \nabla_w \mathcal{L}(x_n, y_n | w_t)$$

$$w_{t+1} := w_t + \varepsilon_t \Delta w_{t+1}$$

$$t := t + 1$$

**RETURN**  $w_n$

## SGD with momentum

**INPUT:**

- \*  $\varepsilon_t$ : controls the speed of convergence
- \*  $\alpha \in (0,1]$ : speed of direction change update
- \* stopping rule

**ALGORITHM:**

initialize  $t = 0$ ,  $w_0$  randomly,  $\Delta w_0 = 0$

**WHILE** stopping rule is not satisfied:

    randomly sample  $I = \{n_1, \dots, n_K\}$  from  $\{1, 2, \dots, N\}$

$$\Delta w_{t+1} = (1 - \alpha)\Delta w_t + \alpha \frac{1}{K} \sum_{n \in I} \nabla_w \mathcal{L}(x_n, y_n | w_t)$$

$$w_{t+1} := w_t + \varepsilon_t \Delta w_{t+1}$$

$$t := t + 1$$

**RETURN**  $w_n$

- Intuition:  $\uparrow$ speed by removing noisy gradients by aggregation over longer history.
- Typically  $\alpha = 0.1$ .

## Other improvements

Other improvements of SGD exist:

- use 2nd order derivative
- Adam, RMSProp, AdaGrad, Adadelata
  - adjust  $\varepsilon_t$  for each dimension individually.
  - important dimensions get  $\downarrow \varepsilon_t$
  - unimportant dimensions get  $\uparrow \varepsilon_t$

# Discussion of SGD

## Advantages

- Simple
- Works online
- A small subset of learning objects may be sufficient for accurate estimation

# Discussion of SGD

## Advantages

- Simple
- Works online
- A small subset of learning objects may be sufficient for accurate estimation

## Drawbacks

- Optimization using 2nd order derivatives converges faster.
- Needs selection of  $\epsilon_t$ :
  - too big: divergence
  - too small: very slow convergence

- If  $\mathcal{L}(\cdot)$  is convex  $\Rightarrow$  convergence to global min from any starting point.
- If  $\mathcal{L}(\cdot)$  is non-convex  $\Rightarrow$  convergence to different local min, depending on starting point.

## Summary

- Gradient descent iteratively optimizes  $L(w)$  in the direction of maximum descent.
  - step takes  $O(N)$
  - $\varepsilon$  should be carefully chosen
- Stochastic gradient descent applies gradient descent to approximation of  $L(w)$ .
  - step takes  $O(K)$
  - requires  $\varepsilon_t \rightarrow 0$  for convergence.
- Feature normalization & momentum speeds up convergence.