

Boosting

Victor Kitov

v.v.kitov@yandex.ru

Linear ensembles

Linear ensemble:

$$F_M(x) = f_0(x) - c_1 f_1(x) + \dots - c_M f_M(x)$$

Regression: $\hat{y}(x) = F_M(x)$

Binary classification: $\text{score}(y|x) = F_M(x)$, $\hat{y}(x) = \text{sign } F_M(x)$

- Notation: $f_1(x), \dots, f_M(x)$ are called *base learners*, *weak learners*, *base models*.
- Too expensive to optimize $f_0(x), f_1(x), \dots, f_M(x)$ and c_1, \dots, c_M jointly for large M .
- Idea: optimize $f_0(x)$ and then each pair $(f_m(x), c_m)$ step-by-step.

Forward stagewise additive modeling (FSAM)



Input:

- training dataset (x_n, y_n) , $n = 1, 2, \dots, N$
- loss function $\mathcal{L}(f, y)$
- parametric form of base learner $f_\theta(x)$
- the number of base learners M .

Output: approximation function $F_M(x) = f_0(x) - \sum_{m=1}^M c_m f_m(x)$


Bagging and boosting

Bagging and boosting




Analogy with playing golf

Analogy with playing golf



Analogy with Taylor expansion

Analogy with Taylor expansion



Forward stagewise additive modeling (FSAM)

- ❶ Fit initial approximation $f_0(x) = \arg \min_f \sum_{n=1}^N \mathcal{L}(f(x_n), y_n)$
- ❷ For $m = 1, 2, \dots, M$:
 - find next best classifier


$$(c_m, f_m) := \arg \min_{f, c} \sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) - cf(x_n), y_n)$$

- reevaluate ensemble

$$F_m(x) := F_{m-1}(x) - c_m f_m(x)$$

Dependency on M

Boosting overfits for high M :



Comments

- M should be determined by performance on validation set.
- Each step should be coarse to leave room for future base learners improvement:
 - initial approximation may be zero or constant
 - optimization can be coarse (just few steps)
 - base learner should be simple
 - such as trees of depth=1,2,3.
- For some loss functions (see Adaboost) we can solve minimization explicitly.
- For general loss functions gradient boosting should be used.

Table of Contents

1 Adaboost

2 Gradient boosting

3 Extensions

Adaboost (discrete version)

Assumptions:

- binary classification task $y \in \{+1, -1\}$
- $f_m(x) \in \{+1, -1\}$, trainable in weighted dataset.
- classification is performed with
$$\hat{y} = \text{sign}\{f_0(x) + c_1 f_1(x) + \dots + c_M f_M(x)\}$$
- optimized loss is $\mathcal{L}(F(x), y) = e^{-yF(x)}$

Optimization in FSAM can be solved explicitly!

Adaboost (discrete version): algorithm

Input:

- training dataset (x_n, y_n) , $n = 1, 2, \dots, N$
- number of additive weak classifiers M
- a family of weak classifiers $f_m(x) \in \{+1, -1\}$
 - should be trainable on weighted datasets.

Output: composite classifier $F_M(x) = \text{sign} \left(\sum_{m=1}^M c_m f_m(x) \right)$

Adaboost (discrete version): algorithm

- ❶ Initialize observation weights $w_n = 1/N$, $n = 1, 2, \dots, N$.
- ❷ for $m = 1, 2, \dots, M$:
 - ❶ fit $f_m(x)$ to training data using weights w_n
 - ❷ compute weighted misclassification rate:

$$E_m = \frac{\sum_{n=1}^N w_n \mathbb{I}[f_m(x_n) \neq y_n]}{\sum_{n=1}^N w_n}$$

- ❸ if $E_m > 0.5$ or $E_m = 0$: terminate procedure.
- ❹ compute $c_m = \frac{1}{2} \ln((1 - E_m)/E_m)$ $E_m < 0.5 \Rightarrow c_m > 0$
- ❺ increase all weights, where misclassification with $f_m(x)$ was made:

$$w_n \leftarrow w_n e^{2c_m} = w_n \left(\frac{1 - E_m}{E_m} \right), \text{ for } n : f_m(x_n) \neq y_n$$

Adaboost derivation

Set $F_0(x) \equiv 0$.

Apply FSAM for $m = 1, 2, \dots, M$:

$$\begin{aligned}(c_m, f_m) &= \arg \min_{c_m, f_m} \sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) + c_m f_m(x_n), y_n) \\&= \arg \min_{c_m, f_m} \sum_{n=1}^N e^{-y_n F_{m-1}(x_n)} e^{-c_m y_n f_m(x_n)} \\&= \arg \min_{c_m, f_m} \sum_{i=1}^N w_n^m e^{-c_m y_n f_m(x_n)}, \quad w_n^m := e^{-y_n F_{m-1}(x_n)}\end{aligned}$$

Adaboost derivation

$$\begin{aligned}
 \sum_{n=1}^N w_n^m e^{-c_m y_n f_m(x_n)} &= \sum_{n: f_m(x_n)=y_n} w_n^m e^{-c_m} + \sum_{n: f_m(x_n) \neq y_n} w_n^m e^{c_m} \\
 &= e^{-c_m} \sum_{n: f_m(x_n)=y_n} w_n^m + e^{c_m} \sum_{n: f_m(x_n) \neq y_n} w_n^m \\
 &= e^{c_m} \sum_{n: f_m(x_n) \neq y_n} w_n^m + e^{-c_m} \sum_{n=1}^N w_n^m - e^{-c_m} \sum_{n: f_m(x_n) \neq y_n} w_n^m \\
 &= e^{-c_m} \sum_n w_n^m + (e^{c_m} - e^{-c_m}) \sum_{n: f_m(x_n) \neq y_n} w_n^m
 \end{aligned}$$

Since $c_m \geq 0$, $f_m(\cdot)$ should be found from

$$f_m(\cdot) = \arg \min_f \sum_{n=1}^N w_n^m \mathbb{I}[f(x_n) \neq y_n]$$

Adaboost derivation

Denote $G(c_m) = \sum_{n=1}^N w_n^m \exp(-c_m y_n f_m(x_n))$. Then

$$\frac{\partial G(c_m)}{\partial c_m} = - \sum_{n=1}^N w_n^m e^{-c_m y_n f_m(x_n)} y_n f_m(x_n) = 0$$

$$- \sum_{n: f_m(x_n)=y_n} w_n^m e^{-c_m} + \sum_{n: f_m(x_n) \neq y_n} w_n^m e^{c_m} = 0$$

$$e^{2c_m} = \frac{\sum_{n: f_m(x_n)=y_n} w_n^m}{\sum_{n: f_m(x_n) \neq y_n} w_n^m}$$

$$c_m = \frac{1}{2} \ln \frac{\left(\sum_{n: f_m(x_n)=y_n} w_n^m \right) / \left(\sum_{n=1}^N w_n^m \right)}{\left(\sum_{n: f_m(x_n) \neq y_n} w_n^m \right) / \left(\sum_{n=1}^N w_n^m \right)} = \frac{1}{2} \ln \frac{1 - E_m}{E_m},$$

$$\text{where } E_m := \frac{\sum_{n=1}^N w_n^m \mathbb{I}[f_m(x_n) \neq y_n]}{\sum_{n=1}^N w_n^m}$$

Adaboost derivation

Weights recalculation:

$$w_n^{m+1} \stackrel{\text{def}}{=} e^{-y_n F_m(x_n)} = e^{-y_n F_{m-1}(x_n)} e^{-y_n c_m f_m(x_n)}$$

Noting that $-y_n f_m(x_n) = 2\mathbb{I}[f_m(x_n) \neq y_n] - 1$, we can rewrite:

$$\begin{aligned} w_n^{m+1} &= e^{-y_n F_{m-1}(x_n)} e^{c_m(2\mathbb{I}[f_m(x_n) \neq y_n] - 1)} = \\ &= w_n^m e^{2c_m \mathbb{I}[f_m(x_n) \neq y_n]} e^{-c_m} \propto w_n^m e^{2c_m \mathbb{I}[f_m(x_n) \neq y_n]} \end{aligned}$$

Comments:

- We can remove common constants from weights.
- $w_n^{m+1} = w_n^m$ for correctly classified objects by $f_m(x)$.
- $w_n^{m+1} = w_n^m e^{2c_m}$ for incorrectly classified objects by $f_m(x)$.
 - so later classifiers will pay more attention to them

Table of Contents

- 1 Adaboost
- 2 Gradient boosting**
- 3 Extensions

Motivation

- Problem: For general loss function L FSAM cannot be solved explicitly
- Analogy with function minimization: when we can't find optimum explicitly we use numerical methods
- Gradient boosting: numerical method for iterative loss minimization

Local linear approximation

Linear approximation \mathcal{L} with $g(x) = \left. \frac{\partial \mathcal{L}(G, y)}{\partial G} \right|_{G=G(x)}$:

$$\mathcal{L}(G(x) - f(x), y) \approx \mathcal{L}(G(x), y) - g(x)f(x)$$

$$\begin{aligned} & \arg \min_{f(x)} \sum_{n=1}^N \mathcal{L}(G(x_n) - f(x_n), y_n) \\ & \approx \arg \min_{f(x)} \sum_{n=1}^N \mathcal{L}(G(x_n), y_n) - g(x_n)f(x_n) \\ & = \arg \min_{f(x)} \sum_{n=1}^N -g(x_n)f(x_n) = \arg \max_{f(x)} \sum_{n=1}^N g(x_n)f(x_n) \end{aligned}$$

$\Rightarrow f(x)$ should approximate $g(x)$, because

$$\arg \max_{f: \|f\| \leq \|g\|} \langle f, g \rangle = g$$

Example: regression

$$\sum_{n=1}^N \left(f_m(x_n) - \frac{\partial \mathcal{L}(G, y)}{\partial G} \Big|_{G=G_{m-1}(x_n)} \right)^2 \rightarrow \min_{f_m}$$

$$\mathcal{L} = \frac{1}{2} (G - y)^2 : f(x) \approx \frac{\partial \mathcal{L}(G, y)}{\partial G} = G - y$$

Example: regression

$$\sum_{n=1}^N \left(f_m(x_n) - \frac{\partial \mathcal{L}(G, y)}{\partial G} \Big|_{G=G_{m-1}(x_n)} \right)^2 \rightarrow \min_{f_m}$$

$$\mathcal{L} = \frac{1}{2} (G - y)^2 : f(x) \approx \frac{\partial \mathcal{L}(G, y)}{\partial G} = G - y$$

$$G_m(x_n) := G_{m-1}(x_n) - c_m f(x) \approx G_{m-1}(x_n) + c_m (y_n - G_{m-1}(x_n))$$

Example: classification

$$\sum_{n=1}^N \left(f_m(x_n) - \frac{\partial \mathcal{L}(G, y)}{\partial G} \Big|_{G=G_{m-1}(x_n)} \right)^2 \rightarrow \min_{f_m}$$

$$\mathcal{L} = [-Gy]_+ : f(x) \approx \frac{\partial \mathcal{L}(G, y)}{\partial G} = \begin{cases} -y, & Gy < 0 \\ 0, & Gy \geq 0 \end{cases}$$

Example: classification

$$\sum_{n=1}^N \left(f_m(x_n) - \frac{\partial \mathcal{L}(G, y)}{\partial G} \Big|_{G=G_{m-1}(x_n)} \right)^2 \rightarrow \min_{f_m}$$

$$\mathcal{L} = [-Gy]_+ : f(x) \approx \frac{\partial \mathcal{L}(G, y)}{\partial G} = \begin{cases} -y, & Gy < 0 \\ 0, & Gy \geq 0 \end{cases}$$

$$G_m(x_n) := G_{m-1}(x_n) - c_m f(x) \approx G_{m-1}(x_n) + \begin{cases} c_m y_n, & G(x_n) y_n < 0 \\ 0, & G(x_n) y_n \geq 0 \end{cases}$$

Gradient descent algorithm

$$L(w) \rightarrow \min_w, \quad g(w) = \nabla_w L(w), \quad w \in \mathbb{R}^N$$

Gradient descent:

```
initialize  $w$   
for  $m = 1, 2, \dots M$ :  
     $g(w) = \nabla_w L(w)$   
     $w = w - \varepsilon g(w)$ 
```

Gradient descent with modified step:

```
initialize  $w$   
for  $m = 1, 2, \dots M$ :  
     $g(w) = \nabla_w L(w)$   
     $c^* = \arg \min_{c > 0} L(w - cg(w))$   
     $w = w - c^* \Delta w$ 
```

Gradient boosting intuition

$$L(F) = \sum_{n=1}^N \mathcal{L}(F^n) \rightarrow \min_F \quad F = [F^1, F^2, \dots, F^N]$$

$$\text{Gradient descent:} \quad F := F - c \nabla L(F)$$

$$\text{Pointwise gradient descent:} \quad F^n := F^n - c \nabla L(F) = F^n - c \nabla \mathcal{L}(F^n)$$

We want generalization to new x , so need functional approximation:

$$F(x) := F(x) - cf(x)$$

$$f(x_n) \approx \nabla \mathcal{L}(F(x_n)) \quad n = 1, 2, \dots, N$$

Gradient boosting

- Now consider

$$L(f(x_1), \dots, f(x_N)) = \sum_{n=1}^N \mathcal{L}(f(x_n), y_n) \rightarrow \min_{f(\cdot)}$$

- Gradient descent performs pointwise optimization, but we need generalization, so we optimize in space of functions.
- Gradient boosting = modified gradient descent in function space:
 - find gradients: $g(x_n) = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \big|_{r=f^{m-1}(x_n)}$
 - fit base learner $f_m(x)$ to $\{(x_n, g(x_n))\}_{n=1}^N$

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- 2 For each step $m = 1, 2, \dots, M$:

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ❶ Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- ❷ For each step $m = 1, 2, \dots, M$:
 - ❶ calculate gradients: $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ❶ Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- ❷ For each step $m = 1, 2, \dots, M$:
 - ❶ calculate gradients: $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
 - ❷ fit $f_m(\cdot)$ to $\{(x_n, z_n)\}_{n=1}^N$, for example by solving

$$\sum_{n=1}^N (f_m(x_n) - g_n)^2 \rightarrow \min_{f_m}$$

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- ② For each step $m = 1, 2, \dots, M$:
 - ① calculate gradients: $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \big|_{r=F_{m-1}(x_n)}$
 - ② fit $f_m(\cdot)$ to $\{(x_n, z_n)\}_{n=1}^N$, for example by solving

$$\sum_{n=1}^N (f_m(x_n) - g_n)^2 \rightarrow \min_{f_m}$$

-
-
- ③ solve univariate optimization problem:

$$c_m = \arg \min_{c>0} \sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) - cf_m(x_n), y_n)$$

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- ② For each step $m = 1, 2, \dots, M$:
 - ① calculate gradients: $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
 - ② fit $f_m(\cdot)$ to $\{(x_n, z_n)\}_{n=1}^N$, for example by solving

$$\sum_{n=1}^N (f_m(x_n) - g_n)^2 \rightarrow \min_{f_m}$$

- ③ solve univariate optimization problem:

$$c_m = \arg \min_{c > 0} \sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) - cf_m(x_n), y_n)$$

- ④ set $F_m(x) = F_{m-1}(x) - c_m f_m(x)$

Gradient boosting

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation $f_0(x)$ (might be taken $f_0(x) \equiv 0$)
- ② For each step $m = 1, 2, \dots, M$:
 - ① calculate gradients: $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \big|_{r=F_{m-1}(x_n)}$
 - ② fit $f_m(\cdot)$ to $\{(x_n, z_n)\}_{n=1}^N$, for example by solving

$$\sum_{n=1}^N (f_m(x_n) - g_n)^2 \rightarrow \min_{f_m}$$

- ③ solve univariate optimization problem:

$$c_m = \arg \min_{c > 0} \sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) - cf_m(x_n), y_n)$$

- ④ set $F_m(x) = F_{m-1}(x) - c_m f_m(x)$

Output: approximation function $F_M(x) = f_0(x) - \sum_{m=1}^M c_m f_m(x)$

Gradient boosting: examples

In gradient boosting

$$\sum_{n=1}^N \left(f_m(x_n) - \frac{\partial \mathcal{L}(r, y)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)^2 \rightarrow \min_{f_m}$$

Consider specific cases:

- $\mathcal{L} = \frac{1}{2} (r - y)^2$
- $\mathcal{L} = [-ry]_+$

Table of Contents

1 Adaboost

2 Gradient boosting

3 Extensions

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- 2 For each step $m = 1, 2, \dots, M$:

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation *with constant*:
$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$
- 2 For each step $m = 1, 2, \dots, M$:
 - 1 calculate gradients $g_n = \left. \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \right|_{r=F_{m-1}(x_n)}$

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- 1 Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- 2 For each step $m = 1, 2, \dots, M$:

- 1 calculate gradients $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
- 2 fit regression tree $f_m(\cdot)$ on $\{(x_n, z_n)\}_{n=1}^N$ with some loss function, get leaf regions $\{R_j^m\}_{j=1}^{J_m}$.

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- ② For each step $m = 1, 2, \dots, M$:

- ① calculate gradients $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
- ② fit regression tree $f_m(\cdot)$ on $\{(x_n, z_n)\}_{n=1}^N$ with some loss function, get leaf regions $\{R_j^m\}_{j=1}^{J_m}$.
- ③ for each terminal region R_j^m , $j = 1, 2, \dots, J_m$ solve univariate optimization problem:

$$\gamma_j^m = \arg \min_{\gamma} \sum_{x_n \in R_j^m} \mathcal{L}(F_{m-1}(x_n) - \gamma, y_n)$$

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- ② For each step $m = 1, 2, \dots, M$:

- ① calculate gradients $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
- ② fit regression tree $f_m(\cdot)$ on $\{(x_n, z_n)\}_{n=1}^N$ with some loss function, get leaf regions $\{R_j^m\}_{j=1}^{J_m}$.
- ③ for each terminal region R_j^m , $j = 1, 2, \dots, J_m$ solve univariate optimization problem:

$$\gamma_j^m = \arg \min_{\gamma} \sum_{x_n \in R_j^m} \mathcal{L}(F_{m-1}(x_n) - \gamma, y_n)$$

- ④ update $F_m(x) = F_{m-1}(x) - \sum_{j=1}^{J_m} \gamma_j^m \mathbb{I}[x \in R_j^m]$

Gradient boosting of trees

Input: training dataset (x_n, y_n) , $n = 1, 2, \dots, N$; loss function $\mathcal{L}(f, y)$ and the number M of successive additive approximations.

- ① Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- ② For each step $m = 1, 2, \dots, M$:

- ① calculate gradients $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
- ② fit regression tree $f_m(\cdot)$ on $\{(x_n, z_n)\}_{n=1}^N$ with some loss function, get leaf regions $\{R_j^m\}_{j=1}^{J_m}$.
- ③ for each terminal region R_j^m , $j = 1, 2, \dots, J_m$ solve univariate optimization problem:

$$\gamma_j^m = \arg \min_{\gamma} \sum_{x_n \in R_j^m} \mathcal{L}(F_{m-1}(x_n) - \gamma, y_n)$$

- ④ update $F_m(x) = F_{m-1}(x) - \sum_{j=1}^{J_m} \gamma_j^m \mathbb{I}[x \in R_j^m]$

Output: approximation function $F_M(x)$

Modification of boosting for trees

- Compared to first method of gradient boosting, boosting of regression trees finds additive coefficients individually for each terminal region R_j^m , not globally for the whole classifier $f_m(x)$.
- This is done to increase accuracy: forward stagewise algorithm cannot be applied to find R_j^m , but it can be applied to find γ_j^m , because second task is solvable for arbitrary L .
- Max depth K : interaction between K features
- Max leaves K : interaction between no more than $\leq K - 1$ features
 - usually $2 \leq K \leq 8$
- M controls underfitting-overfitting trade-off and selected using validation set

Shrinkage & subsampling

- Shrinkage of general GB, step (d):

$$F_m(x) = F_{m-1}(x) - \alpha c_m f_m(x)$$

- Comments:

- $\alpha \in (0, 1]$
- $\alpha \downarrow \implies M \uparrow (\alpha M \approx \text{const})$

- Subsampling

- increases speed of fitting
- may increase accuracy (diversity of base learners \uparrow)

Quadratic loss function approximation

$$\begin{aligned}
 \text{Define } g(x) &= \left. \frac{\partial \mathcal{L}(r, y)}{\partial r} \right|_{r=F(x)}, \quad h(x) = \left. \frac{\partial^2 \mathcal{L}(r, y)}{\partial r^2} \right|_{r=F(x)} \\
 \mathcal{L}(F(x) + f(x), y) &\approx \\
 \mathcal{L}(F(x), y) + g(x)f(x) + \frac{1}{2}h(x)(f(x))^2 &= \\
 \frac{1}{2}h(x) \left(f(x) + \frac{g(x)}{h(x)} \right)^2 + \text{const}(f(x))
 \end{aligned}$$

So $f(x)$ should be fitted to $-g(x)/h(x)$ with weight $h(x)$.

- $h(x) \geq 0$ around local minimum.

Case $y \in \{1, 2, \dots, C\}$

One-vs-all, one-vs-one, error-correcting-codes.

Case $y \in \{1, 2, \dots, C\}$

One-vs-all, one-vs-one, error-correcting-codes.

Alternatively can optimize $\mathcal{L}(F(x), y)$ for $F(x) \in \mathbb{R}^C$

- $F(x) = \{p(y = c|x)\}_{c=1}^C$, y - one-hot encoded true class
- $S(F(x), y) = F(x)^T y = p(y = \text{correct class}|x)$ - score on (x, y)
- $g_n = -\frac{\partial \mathcal{L}(r, y)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \in \mathbb{R}^C$
- $\sum_{n=1}^N (f_m(x_n) - g_n)^2 \rightarrow \min_{f_m}$ yields vector C -dim. regression.
- may use quadratic approximation
 - for efficient inverting of $\left(\frac{\partial^2}{\partial r^2} \mathcal{L}(r, y) \Big|_{r=F(x)} \right)$ may use diagonal approximation.

xgBoost

- One of the most popular algorithms on kaggle.
- Uses decision trees as base learners:
 - $f_m \in \{f(x) = w_{q(x)}\}$,
 - T total number of leaves.
 - $q(x)$ maps $x \in \mathbb{R}^D$ to leaf number
 - $w \in \mathbb{R}^T$ predictions for leaves.

xgBoost

- Loss - 2nd order approximation with **with regularization**:

$$\begin{aligned}
 \mathcal{L}(f_m) &= \sum_{n=1}^N \mathcal{L}(F^{(m-1)}(x_n), y_n) \\
 &\approx \sum_{n=1}^N \left[\mathcal{L}(F^{(m-1)}(x_n), y_n) + g_n f_m(x_n) + \frac{1}{2} h_n f_m^2(x_n) \right] \\
 &\quad + \gamma T + \frac{1}{2} \lambda \sum_{t=1}^T w_t^2
 \end{aligned}$$

- Tree impurity function matches original loss $\mathcal{L}(\cdot, \cdot)$.
- Efficiency optimization:
 - feature values may be discretized for speed
 - parallelization over multiple CPU cores and with GPU

Types of boosting

- Loss function \mathcal{L} :
 - $\mathcal{L}(|f(x) - y|)$ - regression
 - $F(y \cdot \text{score}(y = +1|x))$ - binary classification
 - $\mathcal{L}(F(x), y)$ for $F(x), y \in \mathbb{R}^C$ - multiclass classification
- Optimization
 - analytical (Adaboost)
 - gradient based
 - based on quadratic approximation
- Base learners
 - continuous
 - discrete
- Classification
 - binary
 - multiclass
- Extensions: shrinkage, subsampling