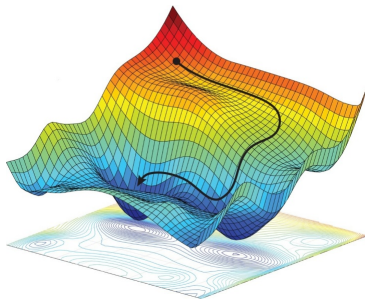


# Neural network optimization

Victor Kitov

[v.v.kitov@yandex.ru](mailto:v.v.kitov@yandex.ru)



# Basic gradient methods

- **Batch gradient descent:** gradient descent using all objects
  - slow for big data
  - not applicable for dynamic data
  - gets stuck in local optima and inflection points (as all gradient based methods)

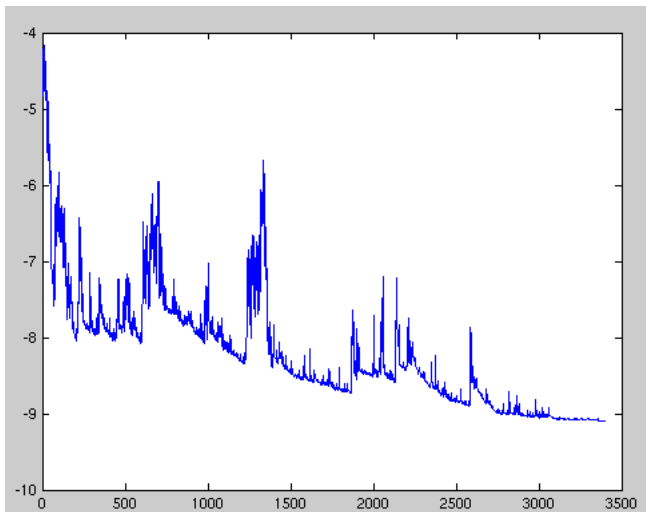
$$w_{t+1} := w_t - \eta \nabla L(\theta; X, Y)$$

- **Stochastic gradient descent:** stochastic descent with sampling one object

$$w_{t+1} := w_t - \eta_t \nabla L(\theta; x_i, y_i)$$

- requires  $\eta_t \rightarrow 0$
- unstable gradient estimate

# SGD convergence example



## Basic gradient methods

- **Minibatch stochastic gradient descent:** stochastic descent with sampling a set of objects

$$w_{t+1} := w_t - \eta_t L(\theta; x_{i+1:i+K}, y_{i+1:i+K})$$

- more accurate gradient estimates
  - faster: computations parallelization over objects in the minibatch
- Difficulties:
  - requires  $\eta_t \rightarrow 0$
  - the same step for different weights
    - better to take less weight, where the function changes sharply.

# Momentum and Nesterov momentum

- **Momentum** ( $\gamma > 0$ ,  $\eta > 0$  - hyperparameters)

$$v_t := \gamma v_{t-1} + \eta \nabla_{\theta} L(\theta)$$

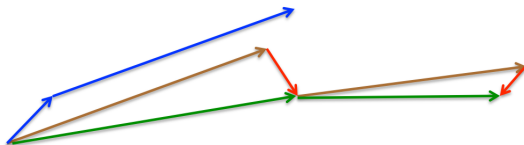
$$w_{t+1} := w_t - v_t$$

- an analogy with a ball rolling down a mountain
- does not stop in small local optima

- **Nesterov Accelerated Gradient** (Nesterov Momentum)

$$v_t := \gamma v_{t-1} + \eta \nabla_{\theta} L(\theta - \gamma v_{t-1})$$

$$w_{t+1} := w_t - v_t$$



# Modifications of SGD

- Denote  $g_t = \nabla L(\theta_t)$ ;  $\theta, g_t \in \mathbb{R}^K$ . Vector operations are elementwise
- **AdaGrad** ( $\varepsilon = 10^{-6}$ )

$$G_t := G_t + g_t^2$$
$$w_{t+1} := w_t - \frac{\eta}{\sqrt{G_t + \varepsilon}} g_t$$

- **RMSprop**

$$E[g^2]_t := \gamma E[g^2]_{t-1} + (1 - \gamma) g_t^2$$
$$w_{t+1} := w_t - \frac{\eta}{\sqrt{E[g^2]_t + \varepsilon}} g_t$$

# Modifications of SGD

- **Adam**=RMSprop+momentum  
( $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$ ,  $\varepsilon = 10^{-8}$ ):

$$m_t := \beta_1 m_{t-1} + (1 - \beta_1) g_1$$

$$v_t := \beta_2 v_{t-1} + (1 - \beta_2) g_1^2$$

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

$$w_{t+1} := w_t - \frac{\eta}{\sqrt{\hat{v}_t} + \varepsilon} \hat{m}_t$$

- **Nadam**: Adam+Nesterov Accelerated Gradient.

# Modifications of SGD

- **AMSGrad**: remembers gradients without exponential forgetting.
  - renormalization  $m_t, v_t$  is not applied and needed.

$$m_t := \beta_1 m_{t-1} + (1 - \beta_1) g_1$$

$$v_t := \beta_2 v_{t-1} + (1 - \beta_2) g_1^2$$

$$\hat{v}_t = \max(\hat{v}_{t-1}, v_t)$$

$$w_{t+1} := w_t - \frac{\eta}{\sqrt{\hat{v}_t} + \varepsilon} \hat{m}_t$$



## Additional improvements<sup>1</sup>

- **Early stopping** - combats overfitting.
- Adding noise to the gradient allows finding the optimum with a larger neighborhood:


$$g_t := g_t + \mathcal{N}(0, \sigma_t^2)$$
$$\sigma_t^2 = \frac{\eta}{(1+t)^\gamma}$$

- **Curriculum learning**
  - first learn on simple objects, then on complex ones
- Improving speed of convergence: batch normalization.

---

<sup>1</sup>More info: <https://ruder.io/deep-learning-optimization-2017/>

# BatchNorm: motivation



- SGD  $w := w - \varepsilon \nabla_w \mathcal{L}(x, y)$  updates all weights on all layers simultaneously.
- Distribution of outputs changes and later layers have to relearn from scratch.
- Also input can shift to saturation region of non-linearity.

# BatchNorm: intro

- Idea: standardize outputs at intermediate layers

$$\tilde{x}_k = \frac{x_k - \mu_k}{\sigma_k}, \quad \mu_k = \mathbb{E}x_k, \sigma_k = \sqrt{\text{Var}(x_k)}$$

- guarantees  $\mathbb{E}\tilde{x}_k = 0$ ,  $\text{Var} \tilde{x}_k = 1$  after weight updates on previous layers.
  - training is faster for later layers
- Training:
  - problem: don't know  $\mu_k, \sigma_k$ 
    - they change dynamically with weight updates
  - solution: estimate them on current minibatch (should be large enough)
- Inference:
  - since now distribution of  $x_k$  is fixed, can estimate  $\mu_k, \sigma_k$  from the whole training set.
  - more efficient: average estimates of  $\mu_k, \sigma_k$  from final minibatches of training.

## BatchNorm: actual version

$$\tilde{x}_k = \alpha_k \frac{x_k - \mu_k}{\sqrt{\sigma_k^2 + \varepsilon}} + \beta_k, \quad \mu_k = \bar{x}_k, \quad \sigma_k = \sqrt{\text{Var}(x_k)}, \quad \varepsilon = 10^{-6}.$$

- Training:
  - $\mu_k, \sigma_k$  estimated from the minibatch
  - $\alpha_k, \beta_k$  - output std.dev. and mean.
    - learned during backpropagation
  - motivation:
    - may cancel normalization effect (e.g. in image->time prediction)
    - flexibility to adjust better to non-linearity
- Inference:
  - $\mu_k, \sigma_k$  estimated fixed to be mean, std.dev. from a wide set of objects.
  - $\alpha_k, \beta_k$  fixed.