

# Word embeddings

Victor Kitov

[v.v.kitov@yandex.ru](mailto:v.v.kitov@yandex.ru)

# Interpretable word embeddings

- $x \in \mathbb{R}^K$ , where  $x^i$  is some  $i$ -th interpretable feature, e.g.
  - $x^1$ : part of speech
  - $x^2$ : gender (for nouns)
  - $x^3$ : tense (for verbs)
  - $x^4$ : starts from capital letter
  - $x^5$ :  $\#$ [letters]
  - $x^6$ : category: machine learning, physics, biology, ...
  - $x^7$ : subcategory: supervised, unsupervised, semi-supervised learning
  - ...
- Need to invent features for each task and extract them.
- Want this to be done automatically!

# Uninterpretable word embeddings

- Clustering words with similar meaning to similar representations.
- **Distributional hypothesis:**  
words have similar meaning  $\Leftrightarrow$  they co-occur together frequently.
- "accuracy of SVM", "SVM gave accuracy", "lower accuracy, compared to SVM"
  - SVM and accuracy are connected!
- Typical dimensionality of embedding  $\in [300, 500]$ .

## Phrase embeddings

We can treat collocations as separate units.

- e.g. fast food, post office, happily married, proud smile.
- may can extract collocations with

$$(w_i, w_j)\text{-collocation} \iff \frac{p(w_i w_j) - \delta}{p(w_i)p(w_j)} > \textit{threshold}$$

$\delta$  - parameter, discouraging rarely co-occurring words as collocations.


# Table of Contents

- 1 Word2vec
- 2 Regularities in embedded space
- 3 Paragraph to vector
- 4 Siamese network

# Word2vec

- For each  $w$  models evaluate:
  - target word embedding  $v_w$
  - context word embedding  $\tilde{v}_w$
- Target&context embeddings may be averaged or concatenated later.

# Continuous bag of words (CBOW)



## Continuous bag of words (CBOW)

CBOW: predict current word given context.


$$\frac{1}{T} \sum_{t=1}^T \ln p(w_t | w_{t-c}, \dots, w_{t-1}, w_{t+1}, \dots, w_{t+c}) \rightarrow \max_{\theta}$$

where  $v_{context} = \sum_{-c \leq i \leq c, i \neq 0} v_{w_{t+i}}$  and

$$p(w_t | w_{t-c}, \dots, w_{t-1}, w_{t+1}, \dots, w_{t+c}) = \frac{\exp(v_{context}^T \tilde{v}_{w_t})}{\sum_{w=1}^V \exp(v_{context}^T \tilde{v}_w)}$$



# Skip-gram model



# Skip-gram model

Skip-gram: predict context, given current word:

$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq i \leq c, i \neq 0} \ln p(w_{t+i} | w_t) \rightarrow \max_{\theta}$$

$$p(w_{t+i} | w_t) = \frac{\exp(v_{w_t}^T \tilde{v}_{w_{t+i}})}{\sum_{w=1}^V \exp(v_{w_t}^T \tilde{v}_w)}$$

## Comments


- May extract embeddings for other objects, appearing in a sequence.
  - symbols, trigrams of symbols (see *FastText*), sentences
  - genes in DNA sequence
  - services ordered by a customer
- May use ensemble of embeddings from different methods.

# Table of Contents

- 1 Word2vec
- 2 Regularities in embedded space
- 3 Paragraph to vector
- 4 Siamese network



## Word forms

Similar in meaning and in spelling words cluster together.



So word embeddings may help in obtaining forms of rare words.


# Word embeddings for different languages are similar<sup>1</sup>



So word embeddings may help in language translation.


<sup>1</sup>Images were manually rotated and scaled.

## Regularities in vector space



So  $(\text{prince} - \text{princess}) + \text{queen} \approx \text{king}$ ! Helps in question answering.

## Regularities in vector space



So  $(\text{Beijing} - \text{China}) + \text{Russia} \approx \text{Moscow}$ ! Helps in question answering.




# Table of Contents

- 1 Word2vec
- 2 Regularities in embedded space
- 3 Paragraph to vector**
- 4 Siamese network

## Paragraph to vector - motivation

- Now need to encode the paragraph (document) into fixed size vector.
- Simple approach: paragraph vector - average of word vectors it contains.
  - or weighted average, considering side information about a word (e.g. stop word/specific term, etc.)
- Alternative: learn paragraph representation.


# Paragraph vector - PV-DM model



## Distributed Memory Model of Paragraph Vectors (PV-DM)

- On training documents are divided into paragraphs. Each paragraph is associated its column in paragraph2vec matrix  $D$ .
- The same task is solved as in CBOW, using vectors and general paragraph information (context).
- Called *Distributed Memory Model of Paragraph Vectors (PV-DM)*.

# Model



- Similar to skip-gram: predicts random sequence of words from the paragraph, using paragraph embedding only.
- Compared to previous, model is simpler: need to estimate only D & softmax parameters; word embedding are not used.
- Called *Distributed Bag of Words version of Paragraph Vector (PV-DBOW)*

# StarSpace<sup>2</sup>

- Facebook library to convert general objects into vector representations.
- Available at [github](#).
- Learns representations of 2 kinds of objects ( $x$  and  $y$ , may coincide) into the same space embedding  $\in \mathbb{R}^D$ .
  - $x \rightarrow a \in \mathbb{R}^D$ ;  $y \rightarrow b \in \mathbb{R}^D$  - can compare different objects!
- Representations are found by sampling randomly matching pair  $(a, b)$  and  $K$  mismatched pairs

$$\sum_{(a,b) \in E^+; b^- \in E^-} L^{batch} (sim(a, b), sim(a, b_1^-), \dots sim(a, b_K^-)) \rightarrow \min_{\{a\}, \{b\}}$$

where  $sim(a, b) = a^T b$  or  $sim(a, b) = \frac{a^T b}{\|a\| \|b\|}$ ;  $L$  - hinge (worked better) or log-loss.

<sup>2</sup>Wu et al. StarSpace: Embed All The Things!

## StarSpace: Applications

- Multiclass classification:  $a$ -feature vector,  $b$ -class.
- Multiclass classification:  $a$ -feature vector,  $b$ -one of matching classes.
- Collaborative filtering:  $a$ -user,  $b$ -item, he likes.
  - does not extend to new users and items
- Collaborative filtering:  $a$ -`"user"`=avg. of items user likes, with excluded item  $b$ .
  - extends to new users, by averaging their item embeddings
  - does not extend to new items
- Collaborative filtering:  $a$ -`"user"`=avg. of items user likes, with excluded item  $b$ . Inside  $a$ =avg(liked items),  $b$ =avg(contained features) - e.g. document consists of words.
  - extends both to new users and new items, as all are featurized.

## StarSpace: Applications

- Link prediction in graphs. Consider graph of relations  $(h, r, t)$  where  $h$ -head concept,  $r$ -relationship,  $t$ -tail concept, e.g. (Beyonce, born in, Houston). Then 2 options possible for sampled  $(h, r, t)$ :
  - $a = (h + r)/2$ ,  $b = t$  or  $a = h$ ,  $b = (r + t)/2$ 
    - thus, can do question answering (Beyonce, born in, ?)
  - Information retrieval:
    - supervised data given:  $a$ -query,  $b$ -relevant document
    - no supervision:  $a$ -random sentence from document  $b$ 
      - both methods produce document embeddings!
  - Word embeddings:  $a$ -surrounding words,  $b$ -central word.
  - Sentence embeddings:  $a, b$ -random sentences from the same document (or close enough to each other)

May find embeddings for 2 tasks simultaneously (sum losses). E.g. jointly learn sentence embeddings and sentence classification according to sentiment (semi-supervised learning).


# Table of Contents

- 1 Word2vec
- 2 Regularities in embedded space
- 3 Paragraph to vector
- 4 Siamese network**



# Siamese network

- Siamese network uses 2 or more duplicate networks, producing embeddings.
- Then these embeddings are compared.

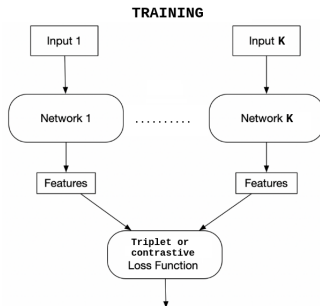


Application: finding similarity or a relationship between two comparable things.

## Application examples

- Classification:
  - input: two objects
  - output: how semantically similar they are (a proxy for class similarity)
- Information retrieval. inputs: are a document and a query,
  - input: a document and a query (may be an image and a visual query-find by image)
  - application: ranking by relevance to a query
- Paraphrase detection:
  - inputs: two sentences
  - output: a score of how similar they are.
- Signature verification
  - inputs: scans of two signatures
  - output: a score that they belong to the same person

# Training



- Loss function:
  - similar objects should have similar embeddings
  - different objects should have distant embeddings



# Losses<sup>4</sup>

## Contrastive loss<sup>3</sup>:

$$\mathbb{I}[y_i = y_j] \|f_{\theta}(x_i) - f_{\theta}(x_j)\|^2 + \mathbb{I}[y_i \neq y_j] \max\{0, \alpha - \|f_{\theta}(x_i) - f_{\theta}(x_j)\|\}^2$$

## Triplet loss:

- $x$  - random object,
- $x, x^+$  are similar (belong to the same class)
- $x, x^-$  are dissimilar (belong to different classes)
- $\alpha > 0$  - hyperparameter (desired margin between positive and negative pairs)

$$\mathcal{L}(x, x^+, x^-) = \max\left\{\|f_{\theta}(x) - f_{\theta}(x^+)\|^2 - \|f_{\theta}(x) - f_{\theta}(x^-)\|^2 - \alpha; 0\right\}$$

---

<sup>3</sup>Chopra et al. Learning a Similarity Metric Discriminatively, with Application to Face Verification.

<sup>4</sup>[Good overview of losses.](#)


## Lifted structural loss



(a) Contrastive embedding



(b) Triplet embedding



(c) Lifted structured embedding

Lifted structural loss<sup>5</sup> utilizes all pairwise similarities.

<sup>5</sup>Deep Metric Learning via Lifted Structured Feature Embedding.

# Lifted structural loss

- Define  $D_{ij} = \|f(x_i) - f(x_j)\|_2$ ,  $P, N$ -the sets of positive and negative pairs.

$$\mathcal{L} = \frac{1}{2|P|} \sum_{(i,j) \in P} \max \left\{ 0; \mathcal{L}_{struct}^{ij} \right\}^2$$

$$\mathcal{L}_{struct}^{ij} = D_{ij} + \max \left\{ \max_{(i,k) \in N} \varepsilon - D_{ik}, \max_{(j,k) \in N} \varepsilon - D_{jk} \right\}$$

- Red parts represent *hard negative mining*.
  - concentrated on hardest negative object, slow training
- Solution: smoothed version:

$$\mathcal{L}_{struct}^{ij} = D_{ij} + \log \left( \sum_{(i,k) \in N} e^{\varepsilon - D_{ik}} + \sum_{(j,k) \in N} e^{\varepsilon - D_{jk}} \right)$$

## Soft-Nearest Neighbors Loss<sup>7</sup>

- Soft-Nearest Neighbors Loss uses many positive objects from the minibatch.

$$\mathcal{L} = -\frac{1}{B} \sum_{i=1}^B \log \left( \frac{\sum_{j=1 \dots B; i \neq j; y_i = y_j} e^{-\|f(x_i), f(x_j)\|^2 / \tau}}{\sum_{k=1, \dots, B; i \neq k} e^{-\|f(x_i), f(x_k)\|^2 / \tau}} \right)$$

- $\tau$  - large  $\Rightarrow$  distance dominated by very similar embeddings, distant embeddings do not contribute.
- See also self-supervised contrastive loss<sup>6</sup>.

Triplet and contrastive losses may be used for metric learning:

$\rho_{\theta}(x, x')$  small for  $x, x'$  belonging to the same class

<sup>6</sup>Khosla et al. Supervised Contrastive Learning

<sup>7</sup>Frosst et al. Analyzing and Improving Representations with the Soft Nearest Neighbor Loss.

## Siamese architecture vs. classification

- Classification learns "what represents each class".
- Siamese network learns "what distinguishes each class from other classes".
- Classification outputs class scores.
- Siamese network outputs distances to each class in embedding space.
- Siamese network
  - is more robust to class imbalance
    - since during training consider instance of each class in turn evenly
    - model learns what makes classes the same/different from other pairs, so few examples of rare class may be enough (*one shot learning*).
  - works well in ensemble with classifier
    - use completely different logic, so much diversity in ensemble
  - requires more training
    - instance based learning=>pairwise learning.



# Embeddings for classifier and Siamese network

Embeddings for classifier (last layer of MLP) and Siamese network for MNIST:

